



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Facultat d'Informàtica de Barcelona



Homework II - Bayesian Modelling

Advanced Statistical Modelling

OCTOBER 23, 2023

Ximena Moure

1 Data Load

First I started by loading the data and plotting it to have a first look at it. I decided to remove the observation with NA values.

```
1 library(R2jags)
2 # Data
3 x <- c(1, 1.5, 2, 2.5, 3, 3.5, 4, 4.5)
4 y <- c(25, 31, 27, 28, 36, 35, NA, 34)
5
6 # Removing NA
7 not_na <- !is.na(y)
8 x <- x[not_na]
9 y <- y[not_na]
10 data_list <- list(x=x, y=y, n=length(y))
11
12 plot(x, y, pch=19, xlab="Fertilizer level", ylab="Yield")
```

Listing 1: Data load and plot




Figure 1: Data visualization

2 Model build

I use JAGS to build the model. The following is a description of it:

- $y[i]$ is the observed data point for the i -th observation.
- The model assumes that each observation $y[i]$ follows a Normal distribution (`dnorm`) with a mean of $b_0 + b_1 \cdot x[i]$ and a precision of τ .
- b_0 and b_1 are the intercept and slope of the linear regression, respectively.
- $x[i]$ is the i -th value of the predictor variable (in this case, the level of fertilizer).
- τ is the precision, which is the reciprocal of the variance ($\tau = \frac{1}{\sigma^2}$).
- b_0 and b_1 are given normal priors with mean 0 and a very small variance (1.0×10^{-7}). This essentially means that before observing the data, the model has a vague or non-informative belief about the values of b_0 and b_1 .
- σ is given a uniform prior between 0 and 100000. This is a very wide range, making this prior quite uninformative as well. It represents a lack of strong prior knowledge about the value of the standard deviation.

The complete code for building the model is depicted in Listing 2

```
1 regression <- "  
2 model {  
3   for (i in 1:n) {  
4     y[i]~ dnorm(b0+b1*x[i], tau)  
5   }  
6   b0 ~ dnorm(0, 1.0E-7)  
7   b1 ~ dnorm(0, 1.0E-7)  
8   sigma ~ dunif(0, 100000)  
9   tau <- pow(sigma, -2)  
10 }  
11 "  
12  
13 ### discarded "Burn", save "Iter", chains "Chain"  
14  
15 Iter <- 10000  
16 Burn <- 1000  
17 Chain <- 2  
18 Thin <- 1  
19  
20 data <- list(y=y, x=x, n=n)  
21  
22 parameters <- c("b0", "b1", "sigma")  
23  
24 potatoes.model <- jags(data, inits=NULL, parameters.to.save=parameters,  
25                        model=textConnection(regression),
```

```

26         n.iter=(Iter*Thin+Burn), n.burnin=Burn, n.thin=Thin, n.chains=Chain
27         )
28 traceplot(potatoes.model, mfrow = c(length(parameters),1), varname = parameters)
29
30 print(potatoes.model)
31
32 attach.jags(potatoes.model)
33 b0 <- b0
34 b1 <- b1
35 sigma <- sigma
36 detach.jags()

```

Listing 2: Model

After running the above code we get the model in figure 2 and the plots depicted in figure 3. From the plots we can see that the process converges.

```

Inference for Bugs model at "3", fit using jags,
  2 chains, each with 11000 iterations (first 1000 discarded)
  n.sims = 20000 iterations saved

```

	mu.vect	sd.vect	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
b0	23.912	4.747	14.696	21.451	23.878	26.379	33.265	1.003	20000
b1	2.700	1.709	-0.660	1.829	2.699	3.582	6.018	1.003	20000
sigma	4.372	2.336	2.054	2.975	3.758	5.030	10.171	1.004	910
deviance	38.146	4.091	33.548	35.157	37.069	39.996	48.756	1.005	800

```

For each parameter, n.eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor (at convergence, Rhat=1).

DIC info (using the rule, pD = var(deviance)/2)
pD = 8.4 and DIC = 46.5
DIC is an estimate of expected predictive error (lower deviance is better).

```

Figure 2: Model result

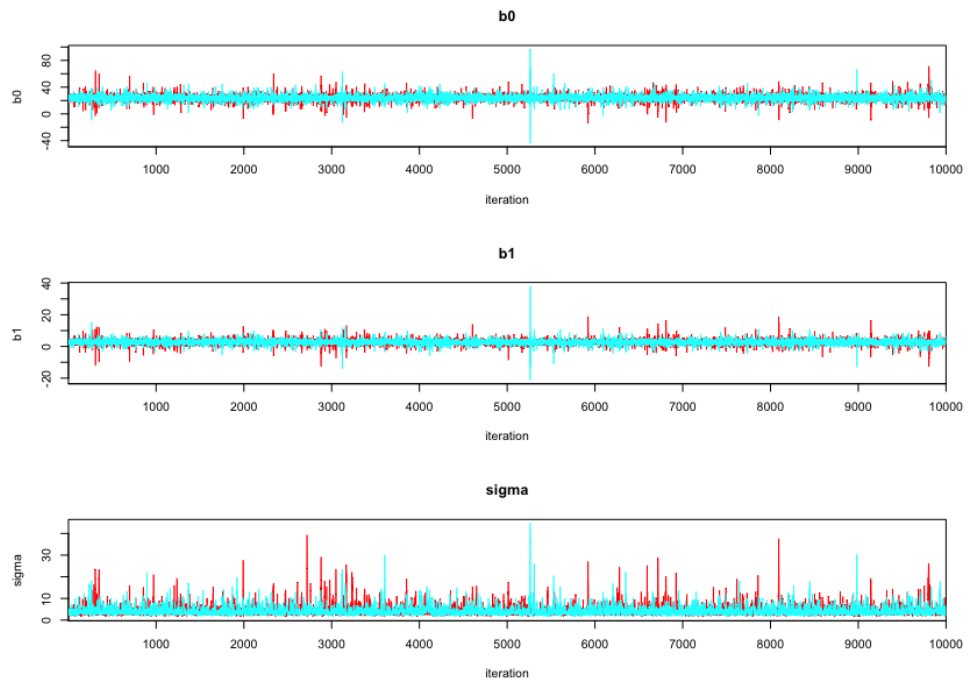


Figure 3: Model plots

3 Draw the posterior distribution for all of them

Now, we can plot the distributions. See Figure 4

```
1 par(mfrow=c(2,2))
2
3 plot(density(b0, adjust = 1.5), main = expression(paste(pi,"(",beta[0],"|y)")), xlab= "
4      " )
5 plot(density(b1, adjust = 1.5), main = expression(paste(pi,"(",beta[1],"|y)")), xlab= "
6      " )
7 plot(density(sigma, adjust = 1.5), main = expression(paste(pi,"(",sigma,"|y)")), xlab=
8      " " )
```

Listing 3: Posterior distribution

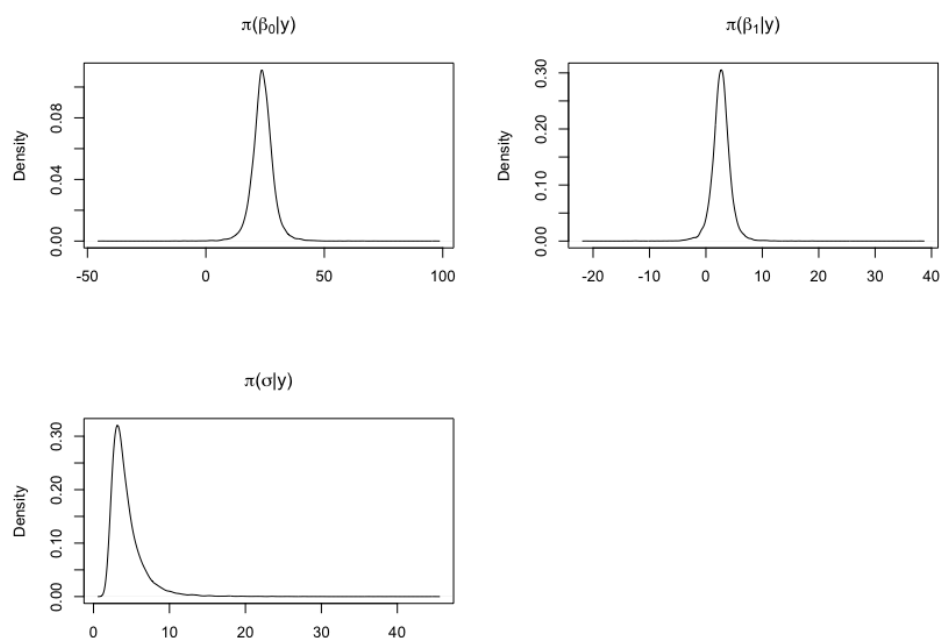


Figure 4: Posterior Distribution

4 Calculate a 95% credible interval for each parameter.

Parameter	Lower Bound	Upper Bound
β_0	14.696	33.265
β_1	-0.661	6.018
σ	2.055	10.171

Table 1: 95% Credible Intervals for each parameter

```
1 # Calculate the 95% credible intervals for b0, b1, and sigma
2 CI_b0 <- quantile(b0, c(0.025, 0.975))
3 CI_b1 <- quantile(b1, c(0.025, 0.975))
4 CI_sigma <- quantile(sigma, c(0.025, 0.975))
5
6 # Print the 95% credible intervals
7 print(CI_b0)
8 print(CI_b1)
9 print(CI_sigma)
```

```
> # Print the 95% credible intervals
> print(CI_b0)
      2.5%      97.5%
14.69641 33.26484
> print(CI_b1)
      2.5%      97.5%
-0.6604218  6.0175296
> print(CI_sigma)
      2.5%      97.5%
 2.05405 10.17069
```

Figure 5: Results credible interval

5 Calculate a 95% credible interval for y given $x = 4$.

To compute a 95% credible interval for y given $x = 4$, we need to generate posterior predictive samples for y at $x = 4$ using the posterior samples of β_0 and β_1 .

Percentile	Lower Bound	Upper Bound
2.5%	22.929	
97.5%		46.120

Table 2: 95% Credible Interval for y given $x = 4$

```
1 # prediction
2
3 M <- length(b0)
4 y.x4 <- rnorm(M, b0+b1*4, sigma)
5
6 par(mfrow=c(2,1))
7 plot(density(y.x4))
8
9
10 quantile(y.x4, c(0.025,0.975))
11
12 summary(y.x4)
```